

# Translation / Transliteration of Vernacular Languages from Signboards

Project ID:21G378383

Review - III

## Group Members

RA1711003030378 Vishnu Teja Chikkala

RA1711003030383 Rajpreet Srivastav

## Supervised By:

Mr. Sunil Kumar

Assistant Professor

Department of Computer Science & Engineering

Faculty of Engineering & Technology

SRM Institute of Science & Technology



# Table of Contents I

- 1 Objective
- 2 Summary of Literature Review
- 3 Architectural Design for Proposed System
- 4 Dataset Specifications
- 5 Methodology / Algorithms / Techniques to be used



# Table of Contents II

6 Partial Implementation

7 Experimental Results

8 References

# Objective I

- India has 22 constitutionally recognized languages written in 13 different scripts. An average traveler, when travelling to a new region, might often get confused with signboards written in an unfamiliar language. It is also impossible to have every signboard in every city / town / village written in 22 different languages, as there will not be enough room to accommodate more than 2 – 3 scripts.
- The objective of this project is to develop a simple and easy-to-use Android mobile app which provides a two-click, picture-to-text, translation / transliteration service for Indian vernacular languages, using deep neural networks and natural language processing models trained for text detection, recognition and translation tasks on collected and freely-available datasets.

## Objective II

- For the scope of this project, we will design a system which works for names (such as road names, city names, shop names, organization names, etc.) which typically are not longer than 4-5 words, and support translation for 1 or 2 languages. Further scope for the project involves including support for more languages and building models to support translation / transliteration of longer pieces of text.

# Summary of Literature Review I

- Indian community faces a “Digital Divide” due to dominance of English as mode of communication in higher education, judiciary, corporate sector and Public administration at Central level whereas the government in states work in their respective regional languages [7]
- India has 22 scheduled languages. While 99 % of the population speak one of these scheduled languages in various dialects (which number in the thousands) [1], according to Census 2011, the total percentage of English speakers is at 10 %, and that too is skewed towards the urban population. [10] Hence, there lies a need for developing NLP architectures for facilitating flow of digital content and information in and between local, national and international levels.

## Summary of Literature Review II

- The above also means that a large percentage of the literate population is either monolingual or bilingual, and across 22 languages, an accessible, easy-to-use and intuitively developed system is required, which enables intercommunication.
- While traditionally NLP has been approached with statistical methods such as Hidden Markov Machines (HMM), Support vector machine(SVM), Conditional Random Field(CRF), Naive Bayes(NB), etc, which take a large amount of tagged/annotated data (corpus) to statistically analyze and learn the language characteristics [3], the research into deep learning or 'connectionist approach' [3] with the use of trained artificial neural networks (ANNs), has gained impetus due to (i) the simplicity of the solution in rapidly prototyping and establishing practically effective systems (ii) the lower cost of annotation of the training data [8], and the fact that they attempt to



# Summary of Literature Review III

more closely emulate the learning process of biological brains, among other reasons. [3], [9], [4]

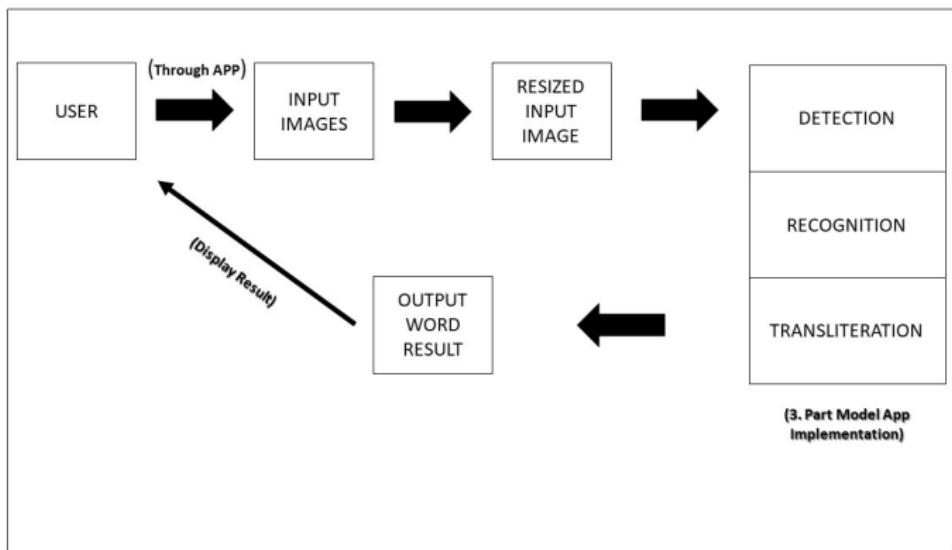
- Particularly, the collection of a uniform corpus and standard datasets for training models remains a challenge across all regional languages. The large number of morphological variations across Indic languages also contributes to this issue.[6], [12]
- Sharma et al., 2017 concluded that almost all existing Indian language machine transliteration systems are based on statistical and hybrid approach [11]
- Kulkarni et al., achieved upto 98% training and 94% testing accuracy on the IndicNLP library for Marathi, using an CNN-LSTM based model architecture. [5]

# Summary of Literature Review IV

- Most of the Indian population accesses digital content through smartphones. In 2019, the number of smartphone users in the country passed 500 million [2], and is estimated to increase to 850 million by 2022 [13]. This, hence, also makes smartphones and smartphone apps in particular an ideal platform on which to launch NLP applications for the wider population, and directly help facilitate flow of information past language barriers.

- The basic functioning of the app is as follows:
  - User captures a photo of the signboard
  - The image is resized so as to be suitable as input to the model
  - The model takes the image as input. The text within the message is detected, extracted and transcribed to target language.
  - The text output (or error message, in case of failure to generate output within threshold confidence), is displayed on screen.

## Architectural Design II



## Figure: Functioning of App



- The artificial neural network (ANN) behind the core functioning of the app is made up of 3 models performing consecutive tasks. That is, the output of a preceding model will be fed as input to the succeeding model, and thus they act as one model unit.

# Architectural Design IV

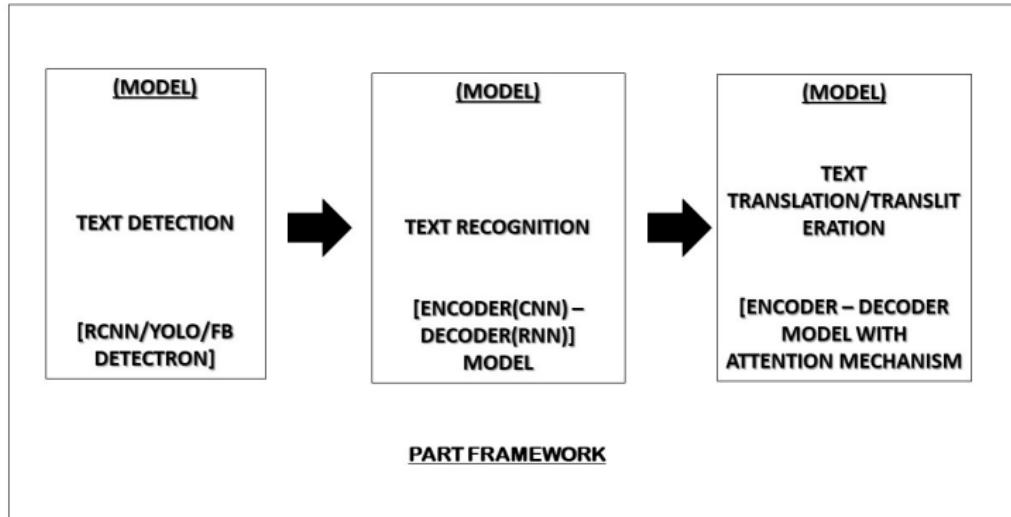


Figure: 3-part model

## Dataset Specification I

- The training dataset for the text detection and text recognition tasks consists of a large number of images containing scene text, synthetically generated. Each of the images has a corresponding annotation, listing the bounding boxes and text script present in the images.



(a) Training Set  
Image

(b) Set  
Annotation

## Dataset Specification II

- The test dataset for text detection and text recognition tasks consists of actual images containing natural scene text.



Figure: Test Set Image

# Dataset Specification III

- The train and test datasets are both in form of xml files containing serialized pairs of source language script and corresponding target language script.

```
1  <?xml version="1.0" encoding="UTF-8"?><TransliterationCorpus CorpusID = "NEWS2012-Training-EnHi-13937" SourceLang = "English" TargetLang = "Hindi" CorpusType = "Training" CorpusSize = "13937" CorpusFormat = "UTF8">
2  <Name ID="1">
3  <SourceName>RAASAVITHAAREE</SourceName>
4  <TargetName ID="1">रासवित्तारी</TargetName>
5  </Name>
6  <Name ID="2">
7  <SourceName>DEOGAN ROAD</SourceName>
8  <TargetName ID="1">देवगन रोड</TargetName>
9  </Name>
10 <Name ID="3">
11 <SourceName>SHATRUMARDAN</SourceName>
12 <TargetName ID="1">शत्रुमरदन</TargetName>
13 </Name>
14 <Name ID="4">
15 <SourceName>MAHIJUBA</SourceName>
16 <TargetName ID="1">महिजुबा</TargetName>
17 </Name>
18 <Name ID="5">
19 <SourceName>SABINE</SourceName>
20 <TargetName ID="1">सैबिन</TargetName>
21 </Name>
22 <Name ID="6">
23 <SourceName>BILL COSBY</SourceName>
24 <TargetName ID="1">बिल कॉस्बी</TargetName>
25 </Name>
26 <Name ID="7">
27 <SourceName>RISHTA KAGAZ KA</SourceName>
28 <TargetName ID="1">रिश्ता कागज़ का</TargetName>
29 </Name>
30 <Name ID="8">
```

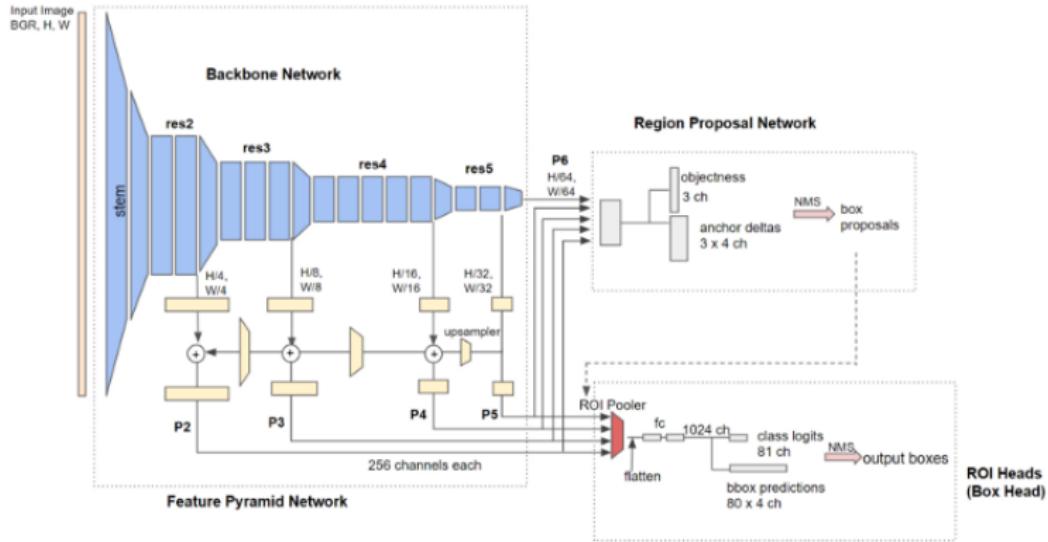
Figure: Transliteration Set



# Methodology / Algorithms / Techniques to be used I

- The text detection task will be carried out by a Faster Region-based Convolutional Network (Faster R-CNN) with Feature Pyramid Network (FPN; for bounding box tightening) from the FAIR Detectron2 kit, which has been trained for object detection on the COCO dataset. We will fine-tune this model to the task at hand by training and validating further on the above scene text database.

## Methodology / Algorithms / Techniques to be used II



**Figure:** Faster R-CNN Object Detection Model

### Methodology / Algorithms / Techniques to be used III

- The text recognition task will be carried out by an encoder-decoder model setup which takes the cropped bounding box of text as input. The encoder is a Convolutional Neural Network (CNN) while the decoder is a Long Short-Term Memory model (LSTM). Connectionist Temporal Classification (CTC) loss will be used to eliminate duplicate recognition of the same letter by adjacent CNN features.

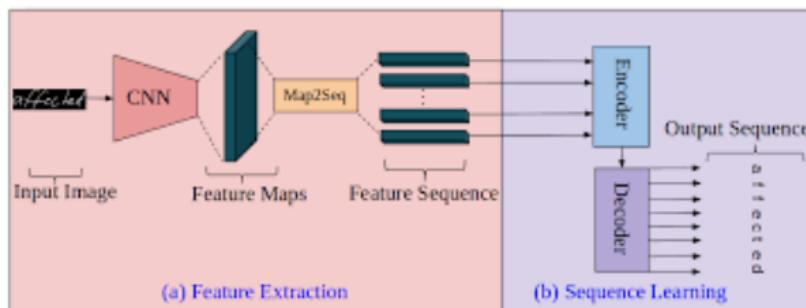


Figure: CNN-LSTM Encoder-Decoder Architecture

## Methodology / Algorithms / Techniques to be used IV

- The transliteration task will be carried by a LSTM - LSTM encoder-decoder model with attention mechanism, which takes source language script as input and generates target language script as output.

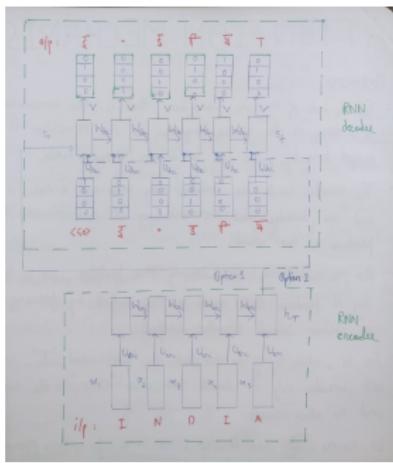


Figure: LSTM-LSTM Encoder-Decoder Architecture

# Methodology / Algorithms / Techniques to be used V

- The app will make use of Google Firebase API to provide UI functionality and user services.
- The model will be mounted on and integrated with the app with the help of Tensorflow Lite.

# Partial Implementation I

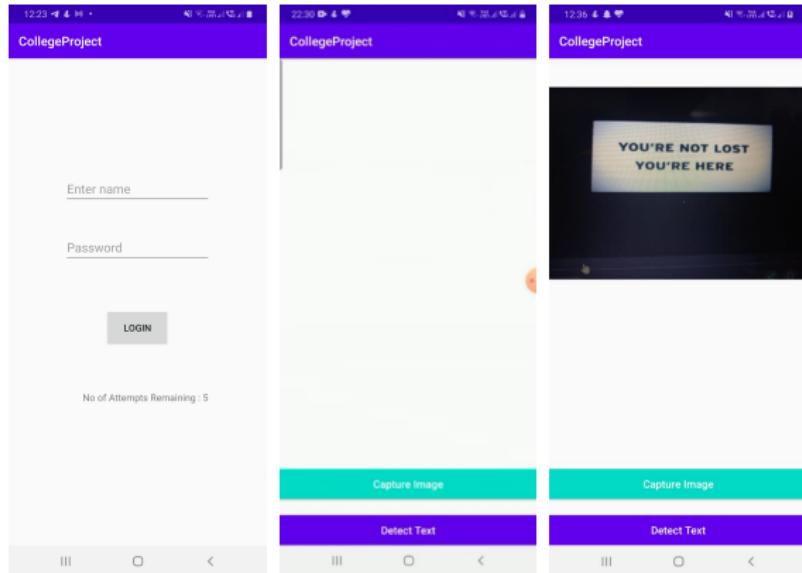


Figure: Skeleton App

# Partial Implementation II

```
def __init__(self, num_chars, rnn_hidden_size=256, dropout=0.1):
    super(CRNN, self).__init__()
    self.num_chars = num_chars
    self.rnn_hidden_size = rnn_hidden_size
    self.dropout = dropout

    # CNN Part 1
    resnet_modules = list(resnet.children())[:-3]
    self.cnn_p1 = nn.Sequential(*resnet_modules)

    # CNN Part 2
    self.cnn_p2 = nn.Sequential(
        nn.Conv2d(256, 256, kernel_size=(3, 6), stride=1, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True)
    )
    self.linear1 = nn.Linear(1024, 256)

    # RNN
    self.rnn1 = nn.GRU(input_size=rnn_hidden_size,
                       hidden_size=rnn_hidden_size,
                       bidirectional=True,
                       batch_first=True)
    self.rnn2 = nn.GRU(input_size=rnn_hidden_size,
                       hidden_size=rnn_hidden_size,
                       bidirectional=True,
                       batch_first=True)
    self.linear2 = nn.Linear(self.rnn_hidden_size*2, num_chars)
```

Figure: Text Recognition Model



# Partial Implementation III

```
MAX_OUTPUT_CHARS = 30
class Transliteration_EncoderDecoder(nn.Module):

    # 'input_size' is the size of the English vocabulary
    # 'output_size' is the size of the Hindi vocabulary
    # 'hidden_size' is a hyper-parameter. More the number of hidden layers in network, greater will be the training accuracy, but also more data will be required
    # 'verbose' enables testing of the model through the execution of it's forward pass
    def __init__(self, input_size, hidden_size, output_size, verbose=False):
        super(Transliteration_EncoderDecoder, self).__init__()

        # Initializes parameters as internal variables for reuse
        self.hidden_size = hidden_size
        self.output_size = output_size

        # Defines the GRU cell for encoder and decoder model
        self.encoder_rnn_cell = nn.GRU(input_size, hidden_size) # Encoder model takes letter by letter input and outputs to hidden layer
        self.decoder_rnn_cell = nn.GRU(output_size, hidden_size) # The input to Decoder model is the output of the previous cell

        self.h2o = nn.Linear(hidden_size, output_size) # In the decoder, converts hidden state to the output ('Hindi' representation)
        self.softmax = nn.LogSoftmax(dim=2) # Softmax layer as this is a classification problem

        self.verbose = verbose

    # Forward pass for the encoder-decoder model
    # The input parameters are the actual input word, the max no. of character in the output word, the device used by the model (CPU or GPU)
    # We can also choose to pass the ground truth (true output) to enable 'teacher forcing'
    def forward(self, input, max_output_chars = MAX_OUTPUT_CHARS, device = 'cpu', ground_truth = None):

        # encoder
        # While the internal classification happens character by character in the model, by passing input as a vector of characters, we get a vectorized output
        hidden = self.encoder_rnn_cell(input)
```

Figure: Text Transliteration Model

# Partial Implementation IV

```
# Hidden state and Output state from previous cell is passed to GRU Cell as input
# The output is the next hidden state and the next Output state
out, decoder_state = self.decoder_rnn_cell(decoder_input, decoder_state)

# Printing the shape of the intermediate output from GRU Cell
if self.verbose:
    print('Decoder intermediate output', out.shape)

out = self.h2o(decoder_state)           # The decoder hidden state is passed through linear layer for classification of character
out = self.softmax(out)                # Softmax layer is applied to produce a probability distribution
outputs.append(out.view(1, -1))         # Output is flattened into a 1D layer

# Printing the shape of this output form
if self.verbose:
    print('Decoder output', out.shape)
    self.verbose = False

# Now, instead of passing the softmax output as is as input to the next GRU Cell iteration, we convert it into a one-hot encoded vector (with the
# This is because of the method of 'Teacher Forcing' which is described below
# Thus, regardless of whether the output of current timestep is from the cell or the ground truth, the input to next timestep will be in one-hot
max_idx = torch.argmax(out, 2, keepdim=True)
if not ground_truth is None:
    max_idx = ground_truth[1].reshape(1, 1, 1)
one_hot = torch.FloatTensor(out.shape).to(device)
one_hot.zero_()
one_hot.scatter_(2, max_idx, 1)

# We do not want gradients to flow through this output-to-input path (only through hidden states), and thus we disable training for this path
decoder_input = one_hot.detach()

return outputs
```

Figure: Text Transliteration Training

# Expected Outcomes I

```
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 3 # NO. of classes = [HINDI, ENGLISH, OTHER]
cfg.TEST.EVAL_PERIOD = 500
os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = CocoTrainer(cfg)
trainer.resume_or_load(resume=False)
trainer.train()
```

[08/01 12:28:05 d2.utils.events]: eta: 0:25:53 iter: 39 total\_loss: 1.523 loss\_cls: 0.602 loss\_box\_reg: 0.758 loss\_rpn\_cls: 0.095 loss\_rpn\_loc: 0.056 time: 1.0614 data\_time
[08/01 12:28:06 d2.utils.events]: eta: 0:25:59 iter: 39 total\_loss: 1.292 loss\_cls: 0.295 loss\_box\_reg: 0.412 loss\_rpn\_cls: 0.024 loss\_rpn\_loc: 0.051 time: 1.0476 data\_time
[08/01 12:29:07 d2.utils.events]: eta: 0:24:48 iter: 79 total\_loss: 1.110 loss\_cls: 0.266 loss\_box\_reg: 0.632 loss\_rpn\_cls: 0.026 loss\_rpn\_loc: 0.051 time: 1.0476 data\_time
[08/01 12:29:18 d2.utils.events]: eta: 0:24:33 iter: 99 total\_loss: 1.007 loss\_cls: 0.251 loss\_box\_reg: 0.478 loss\_rpn\_cls: 0.041 loss\_rpn\_loc: 0.072 time: 1.0517 data\_time
[08/01 12:29:39 d2.utils.events]: eta: 0:24:14 iter: 119 total\_loss: 0.804 loss\_cls: 0.251 loss\_box\_reg: 0.409 loss\_rpn\_cls: 0.029 loss\_rpn\_loc: 0.059 time: 1.0510 data\_time
[08/01 12:30:00 d2.utils.events]: eta: 0:23:48 iter: 139 total\_loss: 0.738 loss\_cls: 0.230 loss\_box\_reg: 0.397 loss\_rpn\_cls: 0.030 loss\_rpn\_loc: 0.053 time: 1.0483 data\_time
[08/01 12:30:21 d2.utils.events]: eta: 0:23:28 iter: 159 total\_loss: 0.698 loss\_cls: 0.229 loss\_box\_reg: 0.381 loss\_rpn\_cls: 0.027 loss\_rpn\_loc: 0.050 time: 1.0484 data\_time
[08/01 12:30:41 d2.utils.events]: eta: 0:23:03 iter: 179 total\_loss: 0.688 loss\_cls: 0.224 loss\_box\_reg: 0.351 loss\_rpn\_cls: 0.024 loss\_rpn\_loc: 0.059 time: 1.0466 data\_time
[08/01 12:31:02 d2.utils.events]: eta: 0:22:43 iter: 199 total\_loss: 0.678 loss\_cls: 0.219 loss\_box\_reg: 0.342 loss\_rpn\_cls: 0.023 loss\_rpn\_loc: 0.051 time: 1.0467 data\_time
[08/01 12:31:23 d2.utils.events]: eta: 0:22:23 iter: 219 total\_loss: 0.631 loss\_cls: 0.214 loss\_box\_reg: 0.350 loss\_rpn\_cls: 0.015 loss\_rpn\_loc: 0.043 time: 1.0467 data\_time
[08/01 12:31:44 d2.utils.events]: eta: 0:22:00 iter: 239 total\_loss: 0.654 loss\_cls: 0.208 loss\_box\_reg: 0.346 loss\_rpn\_cls: 0.017 loss\_rpn\_loc: 0.044 time: 1.0454 data\_time
[08/01 12:32:05 d2.utils.events]: eta: 0:21:38 iter: 259 total\_loss: 0.581 loss\_cls: 0.194 loss\_box\_reg: 0.346 loss\_rpn\_cls: 0.017 loss\_rpn\_loc: 0.044 time: 1.0444 data\_time
[08/01 12:32:25 d2.utils.events]: eta: 0:21:20 iter: 279 total\_loss: 0.613 loss\_cls: 0.184 loss\_box\_reg: 0.355 loss\_rpn\_cls: 0.017 loss\_rpn\_loc: 0.041 time: 1.0439 data\_time
[08/01 12:32:46 d2.utils.events]: eta: 0:20:56 iter: 299 total\_loss: 0.686 loss\_cls: 0.217 loss\_box\_reg: 0.359 loss\_rpn\_cls: 0.025 loss\_rpn\_loc: 0.057 time: 1.0440 data\_time
[08/01 12:33:06 d2.utils.events]: eta: 0:20:39 iter: 319 total\_loss: 0.599 loss\_cls: 0.197 loss\_box\_reg: 0.361 loss\_rpn\_cls: 0.015 loss\_rpn\_loc: 0.043 time: 1.0470 data\_time
[08/01 12:33:27 d2.utils.events]: eta: 0:20:21 iter: 339 total\_loss: 0.581 loss\_cls: 0.196 loss\_box\_reg: 0.362 loss\_rpn\_cls: 0.015 loss\_rpn\_loc: 0.043 time: 1.0470 data\_time
[08/01 12:33:50 d2.utils.events]: eta: 0:19:57 iter: 359 total\_loss: 0.543 loss\_cls: 0.181 loss\_box\_reg: 0.326 loss\_rpn\_cls: 0.010 loss\_rpn\_loc: 0.050 time: 1.0467 data\_time
[08/01 12:34:12 d2.utils.events]: eta: 0:19:38 iter: 379 total\_loss: 0.545 loss\_cls: 0.181 loss\_box\_reg: 0.309 loss\_rpn\_cls: 0.014 loss\_rpn\_loc: 0.047 time: 1.0468 data\_time
[08/01 12:34:31 d2.utils.events]: eta: 0:19:14 iter: 399 total\_loss: 0.620 loss\_cls: 0.163 loss\_box\_reg: 0.295 loss\_rpn\_cls: 0.024 loss\_rpn\_loc: 0.037 time: 1.0448 data\_time
[08/01 12:34:53 d2.utils.events]: eta: 0:18:54 iter: 419 total\_loss: 0.574 loss\_cls: 0.167 loss\_box\_reg: 0.323 loss\_rpn\_cls: 0.017 loss\_rpn\_loc: 0.043 time: 1.0460 data\_time
[08/01 12:35:14 d2.utils.events]: eta: 0:18:34 iter: 439 total\_loss: 0.561 loss\_cls: 0.168 loss\_box\_reg: 0.303 loss\_rpn\_cls: 0.013 loss\_rpn\_loc: 0.040 time: 1.0466 data\_time
[08/01 12:35:34 d2.utils.events]: eta: 0:18:10 iter: 459 total\_loss: 0.543 loss\_cls: 0.171 loss\_box\_reg: 0.332 loss\_rpn\_cls: 0.010 loss\_rpn\_loc: 0.034 time: 1.0440 data\_time
[08/01 12:35:54 d2.utils.events]: eta: 0:17:48 iter: 479 total\_loss: 0.532 loss\_cls: 0.172 loss\_box\_reg: 0.329 loss\_rpn\_cls: 0.009 loss\_rpn\_loc: 0.038 time: 1.0438 data\_time

Figure: Text Detection Training



# Expected Outcomes II

```
[07/28 12:59:29 d2.data.common]: Serializing 27 elements to byte tensors and concatenating them all ...
[07/28 12:59:29 d2.data.common]: Serialized dataset takes 8.81 MB
[07/28 12:59:33 d2.evaluation.evaluator]: Inference done 11/27, 0.393 s / img, ETA=0:00:05
[07/28 12:59:33 d2.evaluation.evaluator]: Inference done 27/27, 0.326 s / img, ETA=0:00:04
[07/28 12:59:33 d2.evaluation.evaluator]: Total inference pure compute time: 0:00:07 (0.324569 s / img per device, on 1 devices)
[07/28 12:59:39 d2.evaluation.evaluator]: Total inference pure compute time: 0:00:07 (0.324569 s / img per device, on 1 devices)
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Preparing results for COCO format ...
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Writing results to ./coco_instances_results.json
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Evaluating predictions ...
loading and preparing results...
Done! (1s)
creating index...
Index created
Running image evaluation...
Evaluate annotation type 'bbox'
2000 (1s-07ms)
Accuracy of evaluation results...
0.98 (1s-07ms).
Average Precision (AP) @ IoU=0.50:0.95 | area= all | methods=100 | = 0.486
Average Precision (AP) @ IoU=0.50:0.95 | area= small | methods=100 | = 0.756
Average Precision (AP) @ IoU=0.50:0.95 | area= medium | methods=100 | = 0.588
Average Precision (AP) @ IoU=0.50:0.95 | area= large | methods=100 | = 0.416
Average Precision (AP) @ IoU=0.50:0.95 | area= tiny | methods=100 | = 0.741
Average Precision (AP) @ IoU=0.50:0.95 | area= small+medium | methods=100 | = 0.454
Average Recall (AR) @ IoU=0.50:0.95 | area= all | methods=100 | = 0.109
Average Recall (AR) @ IoU=0.50:0.95 | area= small | methods=100 | = 0.536
Average Recall (AR) @ IoU=0.50:0.95 | area= medium | methods=100 | = 0.538
Average Recall (AR) @ IoU=0.50:0.95 | area= large | methods=100 | = 0.606
Average Recall (AR) @ IoU=0.50:0.95 | area= tiny | methods=100 | = 0.610
Average Recall (AR) @ IoU=0.50:0.95 | area= small+medium | methods=100 | = 0.519
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
AP | AP50 | AP75 | AP95 | AP99 | AP995
[...]
AP | AP50 | AP75 | AP95 | AP99 | AP995
[...]
AP | AP50 | AP75 | AP95 | AP99 | AP995
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Per category bbox AP:
category | AP | category | AP | category | AP |
[...]
NONE | 57.407 | ENGL150 | 45.207 | OTHER | 43.368
orderedDict([('bbox',
    {'AP': 48.63423889393996,
     'AP50': 45.28907317037224,
     'AP75': 45.28907317037224,
     'AP95': 45.28907317037224,
     'AP99': 45.28907317037224,
     'AP995': 45.28907317037224}])
```

Figure: Text Detection Testing

# Expected Outcomes III

```
ax1.plot(epoch_losses)
```

```
ax1.set_xlabel("Epochs")
```

```
ax1.set_ylabel("Loss")
```

```
ax2.plot(iteration_losses)
```

```
ax2.set_xlabel("Iterations")
```

```
ax2.set_ylabel("Loss")
```

```
plt.show()
```

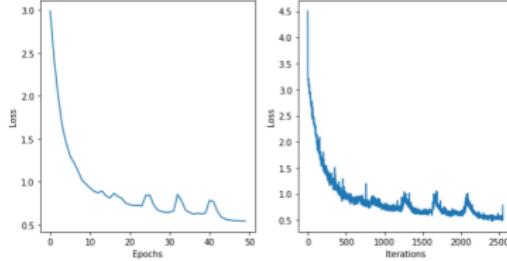


Figure: Text Recognition Training and Testing

# Expected Outcomes IV

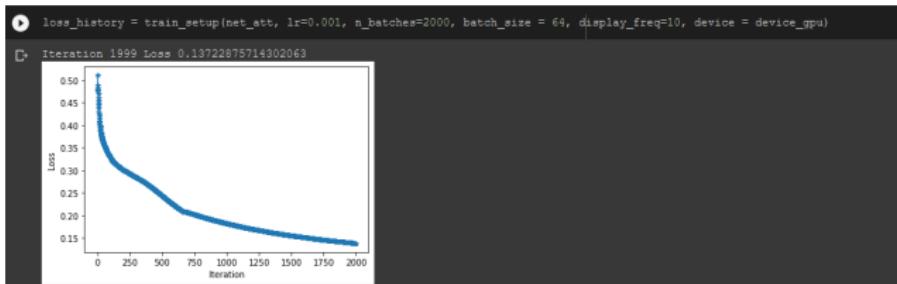


Figure: Text Transliteration Training

```
# Accuracy of the Encoder Decoder model, with and without attention
accuracy = calc_accuracy(net) * 100
accuracy_attn = calc_accuracy(net_att) * 100
print('Accuracy w/o attention ', accuracy)
print('Accuracy with attention', accuracy_attn)

Accuracy w/o attention 69.25025668775679
Accuracy with attention 69.93080114330115
```

Figure: Text Transliteration Testing

# Expected Outcomes V

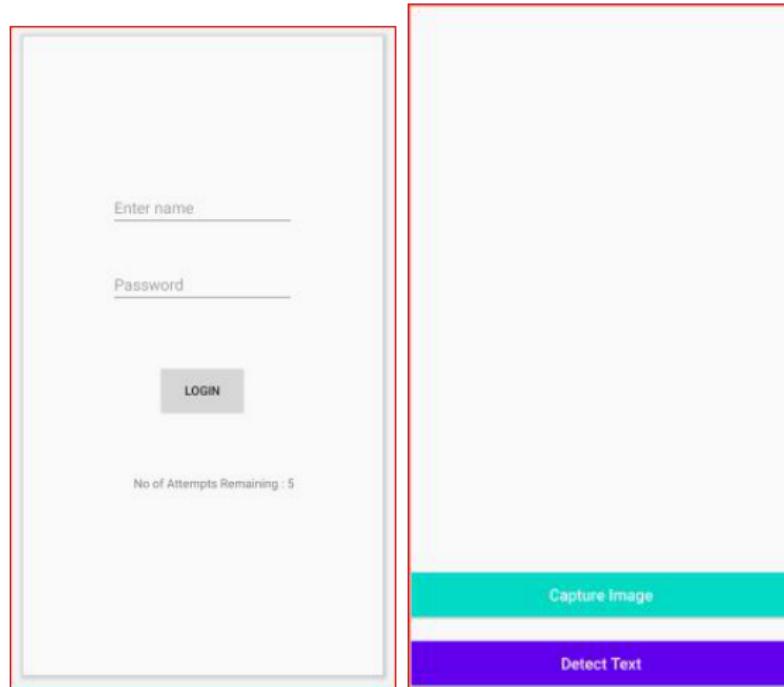


Figure: App Result



# References I

-  More than 19,500 mother tongues spoken in india: Census, Jul 2018.
-  Smartphone users in india crossed 500 million in 2019, states report, Jan 2020.
-  N. P. Desai and V. K. Dabhi.  
Taxonomic survey of hindi language nlp systems.  
*arXiv preprint arXiv:2102.00214*, 2021.
-  T. Deselaers, S. Hasan, O. Bender, and H. Ney.  
A deep learning approach to machine transliteration.  
In *Proceedings of the Fourth Workshop on Statistical Machine Translation*, pages 233–241, 2009.

## References II

-  A. Kulkarni, M. Mandhane, M. Likhitkar, G. Kshirsagar, J. Jagdale, and R. Joshi.  
Experimental evaluation of deep learning models for marathi text classification.  
*arXiv preprint arXiv:2101.04899*, 2021.
-  A. Kunchukuttan, D. Kakwani, S. Golla, A. Bhattacharyya, M. M. Khapra, P. Kumar, et al.  
Ai4bharat-indicnlp corpus: Monolingual corpora and word embeddings for indic languages.  
*arXiv preprint arXiv:2005.00085*, 2020.
-  C. Kurian and K. Kannan Balakrishnan.  
Natural language processing in india prospects and challanges.  
*In Proceedings of the International Conference on "Recent Trends in Computational Science*, 2008.

## References III



J. Philip, V. P. Namboodiri, and C. Jawahar.

A baseline neural machine translation system for indian languages.  
*arXiv preprint arXiv:1907.12437*, 2019.



M. Rosca and T. Breuel.

Sequence-to-sequence neural network models for transliteration.  
*arXiv preprint arXiv:1610.09565*, 2016.



R. S.

In india, who speaks in english, and where?, May 2019.



A. Sharma and D. Rattan.

Machine transliteration for indian languages: A review.

*International Journal of Advanced Research in Computer Science*,  
8(8), 2017.

# References IV



N. Singh.

Nlp for indian languages.  
2020.



[www.ETTelecom.com](http://www.ETTelecom.com).

India to have 820 million smartphone users by 2022 - et telecom, Jul 2020.

# Thank You

