

Translation / Transliteration of Vernacular Languages from Signboards

Project ID:21G378383
Final Review

Group Members
RA1711003030378 Vishnu Teja Chikkala
RA1711003030383 Rajpreet Srivastav

Supervised By:
Mr. Sunil Kumar
Assistant Professor

Department of Computer Science & Engineering
Faculty of Engineering & Technology
SRM Institute of Science & Technology



Table of Contents I

- 1 Objective
- 2 Summary of Literature Review
- 3 Architectural Design for Proposed System
- 4 Dataset Specifications
- 5 Methodology / Algorithms / Techniques to be used



Table of Contents II

6 Partial Implementation

7 Experimental Results

8 Paper Submission Details

9 References



Objective I

- India has 22 constitutionally recognized languages written in 13 different scripts. An average traveler, when travelling to a new region, might often get confused with signboards written in an unfamiliar language. It is also impossible to have every signboard in every city / town / village written in 22 different languages, as there will not be enough room to accommodate more than 2 – 3 scripts.
- The objective of this project is to develop a simple and easy-to-use Android mobile app which provides a two-click, picture-to-text, translation / transliteration service for Indian vernacular languages, using deep neural networks and natural language processing models trained for text detection, recognition and translation tasks on collected and freely-available datasets.

Objective II

- For the scope of this project, we will design a system which works for names (such as road names, city names, shop names, organization names, etc.) which typically are not longer than 4-5 words, and support translation for 1 or 2 languages. Further scope for the project involves including support for more languages and building models to support translation / transliteration of longer pieces of text.

Summary of Literature Review I

- Indian community faces a “Digital Divide” due to dominance of English as mode of communication in higher education, judiciary, corporate sector and Public administration at Central level whereas the government in states work in their respective regional languages [9]
- India has 22 scheduled languages. While 99 % of the population speak one of these scheduled languages in various dialects (which number in the thousands) [1], according to Census 2011, the total percentage of English speakers is at 10 %, and that too is skewed towards the urban population. [12] Hence, there lies a need for developing NLP architectures for facilitating flow of digital content and information in and between local, national and international levels.



Summary of Literature Review II

- The above also means that a large percentage of the literate population is either monolingual or bilingual, and across 22 languages, an accessible, easy-to-use and intuitively developed system is required, which enables intercommunication.
- While traditionally NLP has been approached with statistical methods such as Hidden Markov Machines (HMM), Support vector machine(SVM), Conditional Random Field(CRF), Naive Bayes(NB), etc, which take a large amount of tagged/annotated data (corpus) to statistically analyze and learn the language characteristics [5], the research into deep learning or 'connectionist approach' [5] with the use of trained artificial neural networks (ANNs), has gained impetus due to (i) the simplicity of the solution in rapidly prototyping and establishing practically effective systems (ii) the lower cost of annotation of the training data [10], and the fact that they attempt



Summary of Literature Review III

to more closely emulate the learning process of biological brains, among other reasons. [5], [11], [6]

- Particularly, the collection of a uniform corpus and standard datasets for training models remains a challenge across all regional languages. The large number of morphological variations across Indic languages also contributes to this issue.[8], [14]
- Sharma et al., 2017 concluded that almost all existing Indian language machine transliteration systems are based on statistical and hybrid approach [13]

Summary of Literature Review IV

- Most of the Indian population accesses digital content through smartphones. In 2019, the number of smartphone users in the country passed 500 million [2], and is estimated to increase to 850 million by 2022 [15]. This, hence, also makes smartphones and smartphone apps in particular an ideal platform on which to launch NLP applications for the wider population, and directly help facilitate flow of information past language barriers.

Summaries of a few studies and models proposed in this field are as follows:

Authors	Model	Accuracy
Bhanja et al., 2019 [4]	CNN-LSTM (ResNet)	93% (NITS-LD), 89% (OGI-MLT)
Kulkarni et al., 2021 [7]	CNN-LSTM	94% (IndicNLP)
J. Philip et al, 2019 [10]	IL-Multi	40% (untrained; transferred learning)
Arafat, S. Y. et al, 2020 [3]	FasterR-CNN - CNN - RRNN	99%

- The basic functioning of the app is as follows:
 - User captures a photo of the signboard
 - The image is resized so as to be suitable as input to the model
 - The model takes the image as input. The text within the message is detected, extracted and transcribed to target language.
 - The text output (or error message, in case of failure to generate output within threshold confidence), is displayed on screen.

Architectural Design II

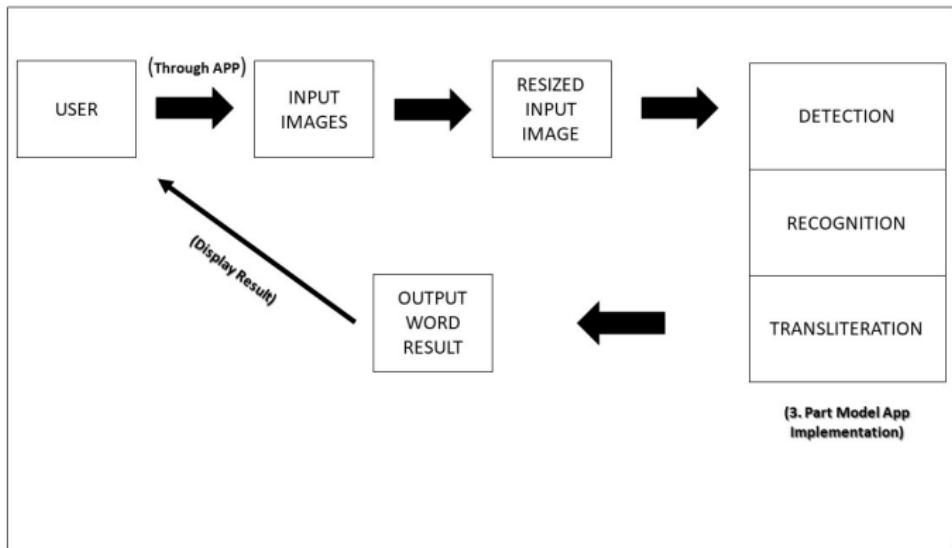


Figure: Functioning of App



Architectural Design III

- The artificial neural network (ANN) behind the core functioning of the app is made up of 3 models performing consecutive tasks. That is, the output of a preceding model will be fed as input to the succeeding model, and thus they act as one model unit.



Architectural Design IV

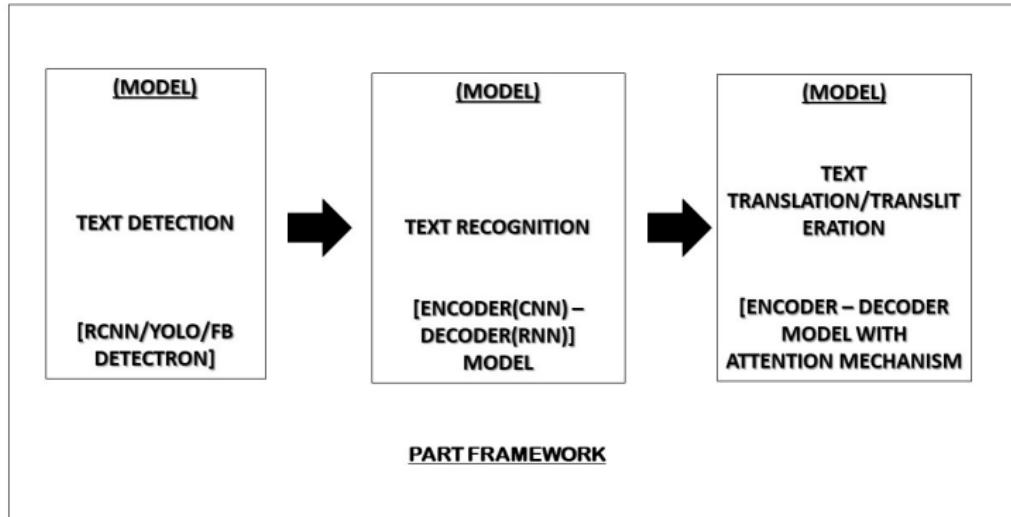


Figure: 3-part model

Dataset Specification I

- The training dataset for the text detection and text recognition tasks consists of a large number of images containing scene text, synthetically generated. Each of the images has a corresponding annotation, listing the bounding boxes and text script present in the images.



(a) Training Set
Image

(b) Set
Annotation

```
3 3252099 79 479644 91.47773 27.37761 111.44640 87.39837 129.44270 144.3673 477  
185.26473 165.49769 164.67719 195.45846 296.7277 296.54646 411.4902 412.1232 1796  
305.44893 305.44893 305.44893 305.44893 305.44893 305.44893 305.44893 305.44893 305.44893 305.44893  
275.59427 311.40835 302.344 285.49373 134.48113 136.5007 175.5007 176.5007 146.49574 19  
30.48979 140.09719 144.093 37.49711 162.09381 181.17182 274.17178 275.04119 19  
163.05446 164.05446 165.05446 166.05446 167.05446 168.05446 169.05446 170.05446 171.05446 172.05446  
225.05446 226.05446 227.05446 228.05446 229.05446 230.05446 231.05446 232.05446 233.05446 234.05446  
345.05446 312.05230 312.05446 344.079 344.079 344.079 344.079 344.079 344.079 344.079  
107.42642 164.81768 165.15936 167.48462 241.99946 242.49946 291.64646 292.89746 19996  
107.42642 164.81768 165.15936 167.48462 241.99946 242.49946 291.64646 292.89746 19996
```

Dataset Specification II

- The test dataset for text detection and text recognition tasks consists of actual images containing natural scene text.



Figure: Test Set Image

Dataset Specification III

- The train and test datasets are both in form of xml files containing serialized pairs of source language script and corresponding target language script.

```
1  <?xml version="1.0" encoding="UTF-8"?><TransliterationCorpus CorpusID = "NEWS2012-Training-EnHi-13937" SourceLang = "English" TargetLang = "Hindi" CorpusType = "Training" CorpusSize = "13937" CorpusFormat = "UTF8">
2  <Name ID="1">
3  <SourceName>RAASAVITHAAREE</SourceName>
4  <TargetName ID="1">रासवित्तारी</TargetName>
5  </Name>
6  <Name ID="2">
7  <SourceName>DEOGAN ROAD</SourceName>
8  <TargetName ID="1">देवगन रोड</TargetName>
9  </Name>
10 <Name ID="3">
11 <SourceName>SHATRUMARDAN</SourceName>
12 <TargetName ID="1">शत्रुमरदन</TargetName>
13 </Name>
14 <Name ID="4">
15 <SourceName>MAHIJUBA</SourceName>
16 <TargetName ID="1">महिजुबा</TargetName>
17 </Name>
18 <Name ID="5">
19 <SourceName>SABINE</SourceName>
20 <TargetName ID="1">सैबिन्</TargetName>
21 </Name>
22 <Name ID="6">
23 <SourceName>BILL COSBY</SourceName>
24 <TargetName ID="1">बिल कॉस्बी</TargetName>
25 </Name>
26 <Name ID="7">
27 <SourceName>RISHTA KAGAZ KA</SourceName>
28 <TargetName ID="1">रिश्ता कागज़ का</TargetName>
29 </Name>
30 <Name ID="8">
```

Figure: Transliteration Set



Methodology / Algorithms / Techniques to be used I

- The text detection task will be carried out by a Faster Region-based Convolutional Network (Faster R-CNN) with Feature Pyramid Network (FPN; for bounding box tightening) from the FAIR Detectron2 kit, which has been trained for object detection on the COCO dataset. We will fine-tune this model to the task at hand by training and validating further on the above scene text database.

Methodology / Algorithms / Techniques to be used II

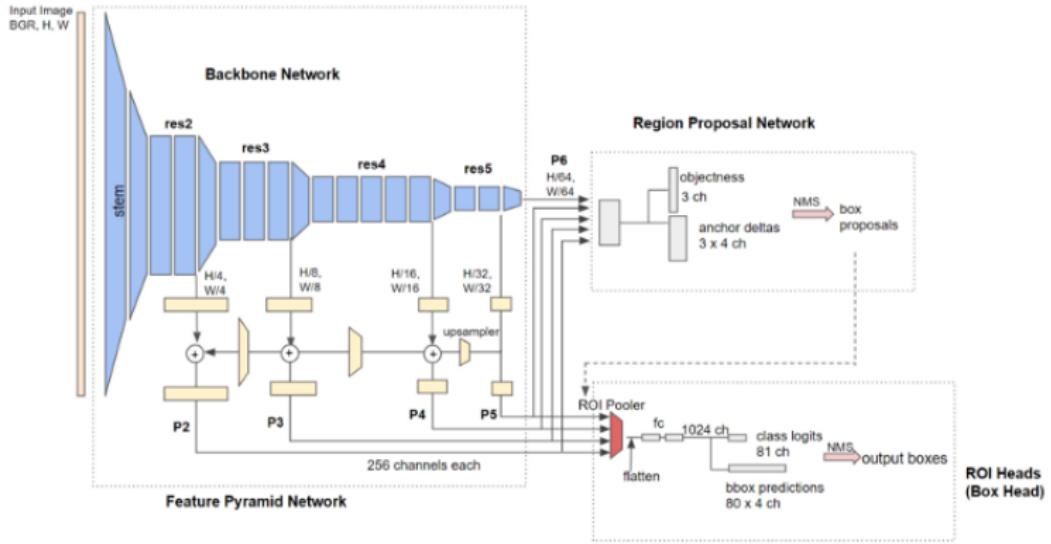


Figure: Faster R-CNN Object Detection Model

Methodology / Algorithms / Techniques to be used III

- The text recognition task will be carried out by an encoder-decoder model setup which takes the cropped bounding box of text as input. The encoder is a Convolutional Neural Network (CNN) while the decoder is a Long Short-Term Memory model (LSTM). Connectionist Temporal Classification (CTC) loss will be used to eliminate duplicate recognition of the same letter by adjacent CNN features.

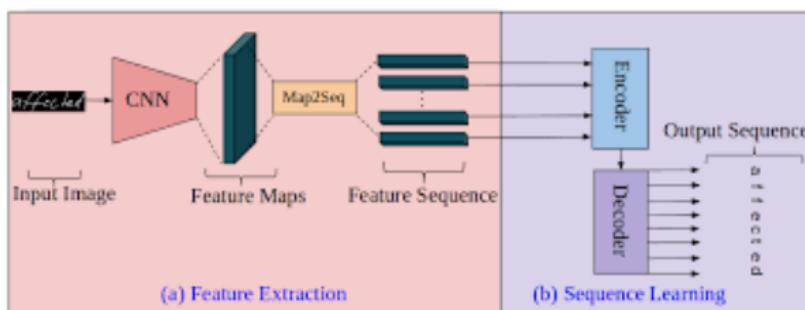


Figure: CNN-LSTM Encoder-Decoder Architecture

Methodology / Algorithms / Techniques to be used IV

- The transliteration task will be carried by a LSTM - LSTM encoder-decoder model with attention mechanism, which takes source language script as input and generates target language script as output.

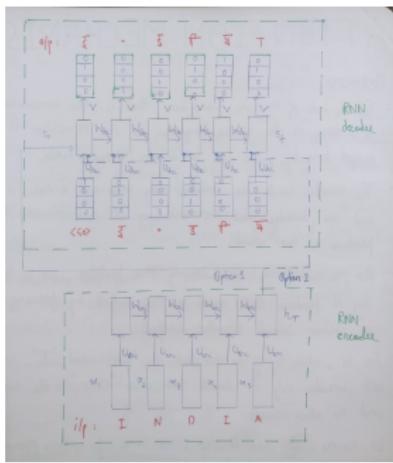


Figure: LSTM-LSTM Encoder-Decoder Architecture

Methodology / Algorithms / Techniques to be used V

- The app will make use of Google Firebase API to provide UI functionality and user services.
- The model will be mounted on and integrated with the app with the help of Tensorflow Lite.

Partial Implementation I

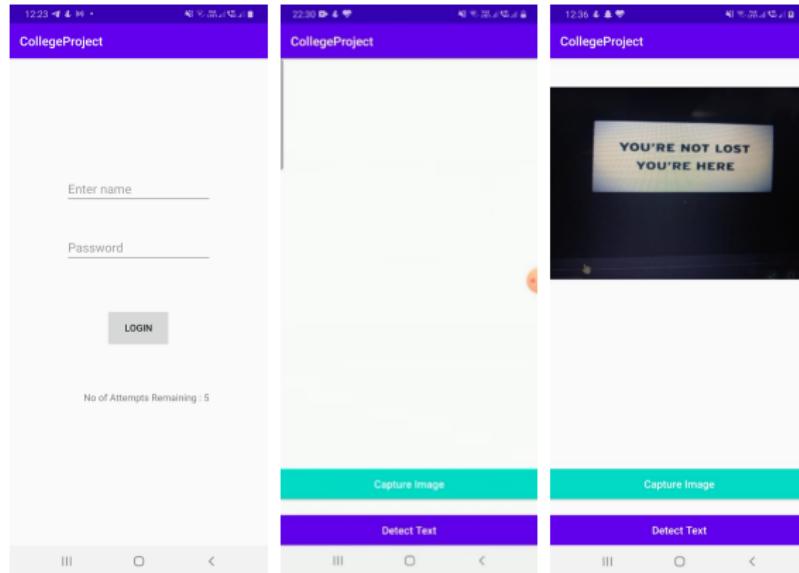


Figure: Skeleton App

Partial Implementation II

```
def __init__(self, num_chars, rnn_hidden_size=256, dropout=0.1):
    super(CRNN, self).__init__()
    self.num_chars = num_chars
    self.rnn_hidden_size = rnn_hidden_size
    self.dropout = dropout

    # CNN Part 1
    resnet_modules = list(resnet.children())[:-3]
    self.cnn_p1 = nn.Sequential(*resnet_modules)

    # CNN Part 2
    self.cnn_p2 = nn.Sequential(
        nn.Conv2d(256, 256, kernel_size=(3, 6), stride=1, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True)
    )
    self.linear1 = nn.Linear(1024, 256)

    # RNN
    self.rnn1 = nn.GRU(input_size=rnn_hidden_size,
                       hidden_size=rnn_hidden_size,
                       bidirectional=True,
                       batch_first=True)
    self.rnn2 = nn.GRU(input_size=rnn_hidden_size,
                       hidden_size=rnn_hidden_size,
                       bidirectional=True,
                       batch_first=True)
    self.linear2 = nn.Linear(self.rnn_hidden_size*2, num_chars)
```

Figure: Text Recognition Model

Partial Implementation III

```
MAX_OUTPUT_CHARS = 30
class Transliteration_EncoderDecoder(nn.Module):

    # 'input_size' is the size of the English vocabulary
    # 'output_size' is the size of the Hindi vocabulary
    # 'hidden_size' is a hyper-parameter. More the number of hidden layers in network, greater will be the training accuracy, but also more data will be required
    # 'verbose' enables testing of the model through the execution of it's forward pass
    def __init__(self, input_size, hidden_size, output_size, verbose=False):
        super(Transliteration_EncoderDecoder, self).__init__()

        # Initializes parameters as internal variables for reuse
        self.hidden_size = hidden_size
        self.output_size = output_size

        # Defines the GRU cell for encoder and decoder model
        self.encoder_rnn_cell = nn.GRU(input_size, hidden_size) # Encoder model takes letter by letter input and outputs to hidden layer
        self.decoder_rnn_cell = nn.GRU(output_size, hidden_size) # The input to Decoder model is the output of the previous cell

        self.h2o = nn.Linear(hidden_size, output_size) # In the decoder, converts hidden state to the output ('Hindi' representation)
        self.softmax = nn.LogSoftmax(dim=2) # Softmax layer as this is a classification problem

        self.verbose = verbose

    # Forward pass for the encoder-decoder model
    # The input parameters are the actual input word, the max no. of character in the output word, the device used by the model (CPU or GPU)
    # We can also choose to pass the ground truth (true output) to enable 'teacher forcing'
    def forward(self, input, max_output_chars = MAX_OUTPUT_CHARS, device = 'cpu', ground_truth = None):

        # encoder
        # While the internal classification happens character by character in the model, by passing input as a vector of characters, we get a vectorized output
        hidden = self.encoder_rnn_cell(input)
```

Figure: Text Transliteration Model

Partial Implementation IV

```
# Hidden state and Output state from previous cell is passed to GRU Cell as input
# The output is the next hidden state and the next Output state
out, decoder_state = self.decoder_rnn_cell(decoder_input, decoder_state)

# Printing the shape of the intermediate output from GRU Cell
if self.verbose:
    print('Decoder intermediate output', out.shape)

out = self.h2o(decoder_state)           # The decoder hidden state is passed through linear layer for classification of character
out = self.softmax(out)                # Softmax layer is applied to produce a probability distribution
outputs.append(out.view(1, -1))         # Output is flattened into a 1D layer

# Printing the shape of this output form
if self.verbose:
    print('Decoder output', out.shape)
    self.verbose = False

# Now, instead of passing the softmax output as is as input to the next GRU Cell iteration, we convert it into a one-hot encoded vector (with the
# This is because of the method of 'Teacher Forcing' which is described below
# Thus, regardless of whether the output of current timestep is from the cell or the ground truth, the input to next timestep will be in one-hot
max_idx = torch.argmax(out, 2, keepdim=True)
if not ground_truth is None:
    max_idx = ground_truth[1].reshape(1, 1, 1)
one_hot = torch.FloatTensor(out.shape).to(device)
one_hot.zero_()
one_hot.scatter_(2, max_idx, 1)

# We do not want gradients to flow through this output-to-input path (only through hidden states), and thus we disable training for this path
decoder_input = one_hot.detach()

return outputs
```

Figure: Text Transliteration Training

Experimental Results I

```
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 3 # NO. of classes = [HINDI, ENGLISH, OTHER]
cfg.TEST.EVAL_PERIOD = 500
os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = CocoTrainer(cfg)
trainer.resume_or_load(resume=False)
trainer.train()
```

[08/01 12:28:51 d2.utils.events]: eta: 0:25:51 iter: 39 total_loss: 1.523 loss_cls: 0.602 loss_box_reg: 0.758 loss_rpn_cls: 0.095 loss_rpn_loc: 0.056 time: 1.0614 data_time
[08/01 12:29:02 d2.utils.events]: eta: 0:25:59 iter: 39 total_loss: 1.295 loss_cls: 0.295 loss_box_reg: 0.424 loss_rpn_cls: 0.026 loss_rpn_loc: 0.051 time: 1.0476 data_time
[08/01 12:29:07 d2.utils.events]: eta: 0:24:48 iter: 79 total_loss: 1.110 loss_cls: 0.266 loss_box_reg: 0.432 loss_rpn_cls: 0.026 loss_rpn_loc: 0.051 time: 1.0476 data_time
[08/01 12:29:18 d2.utils.events]: eta: 0:24:33 iter: 99 total_loss: 1.007 loss_cls: 0.251 loss_box_reg: 0.478 loss_rpn_cls: 0.041 loss_rpn_loc: 0.072 time: 1.0517 data_time
[08/01 12:29:39 d2.utils.events]: eta: 0:24:14 iter: 119 total_loss: 0.804 loss_cls: 0.251 loss_box_reg: 0.409 loss_rpn_cls: 0.029 loss_rpn_loc: 0.059 time: 1.0510 data_time
[08/01 12:30:00 d2.utils.events]: eta: 0:23:48 iter: 139 total_loss: 0.738 loss_cls: 0.230 loss_box_reg: 0.397 loss_rpn_cls: 0.030 loss_rpn_loc: 0.053 time: 1.0483 data_time
[08/01 12:30:21 d2.utils.events]: eta: 0:23:28 iter: 159 total_loss: 0.698 loss_cls: 0.229 loss_box_reg: 0.381 loss_rpn_cls: 0.027 loss_rpn_loc: 0.050 time: 1.0484 data_time
[08/01 12:30:41 d2.utils.events]: eta: 0:23:03 iter: 179 total_loss: 0.688 loss_cls: 0.224 loss_box_reg: 0.351 loss_rpn_cls: 0.024 loss_rpn_loc: 0.049 time: 1.0466 data_time
[08/01 12:30:52 d2.utils.events]: eta: 0:22:53 iter: 199 total_loss: 0.678 loss_cls: 0.221 loss_box_reg: 0.342 loss_rpn_cls: 0.023 loss_rpn_loc: 0.048 time: 1.0467 data_time
[08/01 12:31:23 d2.utils.events]: eta: 0:22:23 iter: 219 total_loss: 0.631 loss_cls: 0.221 loss_box_reg: 0.350 loss_rpn_cls: 0.015 loss_rpn_loc: 0.043 time: 1.0467 data_time
[08/01 12:31:44 d2.utils.events]: eta: 0:22:00 iter: 239 total_loss: 0.654 loss_cls: 0.200 loss_box_reg: 0.346 loss_rpn_cls: 0.017 loss_rpn_loc: 0.044 time: 1.0454 data_time
[08/01 12:32:05 d2.utils.events]: eta: 0:21:38 iter: 259 total_loss: 0.581 loss_cls: 0.194 loss_box_reg: 0.346 loss_rpn_cls: 0.017 loss_rpn_loc: 0.044 time: 1.0444 data_time
[08/01 12:32:25 d2.utils.events]: eta: 0:21:20 iter: 279 total_loss: 0.613 loss_cls: 0.184 loss_box_reg: 0.355 loss_rpn_cls: 0.017 loss_rpn_loc: 0.041 time: 1.0439 data_time
[08/01 12:32:46 d2.utils.events]: eta: 0:20:56 iter: 299 total_loss: 0.686 loss_cls: 0.217 loss_box_reg: 0.359 loss_rpn_cls: 0.025 loss_rpn_loc: 0.057 time: 1.0440 data_time
[08/01 12:33:08 d2.utils.events]: eta: 0:20:39 iter: 319 total_loss: 0.599 loss_cls: 0.197 loss_box_reg: 0.361 loss_rpn_cls: 0.015 loss_rpn_loc: 0.043 time: 1.0470 data_time
[08/01 12:33:29 d2.utils.events]: eta: 0:20:21 iter: 339 total_loss: 0.581 loss_cls: 0.196 loss_box_reg: 0.362 loss_rpn_cls: 0.015 loss_rpn_loc: 0.043 time: 1.0470 data_time
[08/01 12:33:50 d2.utils.events]: eta: 0:19:57 iter: 359 total_loss: 0.543 loss_cls: 0.181 loss_box_reg: 0.326 loss_rpn_cls: 0.010 loss_rpn_loc: 0.050 time: 1.0467 data_time
[08/01 12:34:12 d2.utils.events]: eta: 0:19:36 iter: 379 total_loss: 0.545 loss_cls: 0.181 loss_box_reg: 0.309 loss_rpn_cls: 0.024 loss_rpn_loc: 0.047 time: 1.0468 data_time
[08/01 12:34:31 d2.utils.events]: eta: 0:19:14 iter: 399 total_loss: 0.620 loss_cls: 0.163 loss_box_reg: 0.395 loss_rpn_cls: 0.016 loss_rpn_loc: 0.056 time: 1.0467 data_time
[08/01 12:34:53 d2.utils.events]: eta: 0:18:54 iter: 419 total_loss: 0.574 loss_cls: 0.167 loss_box_reg: 0.326 loss_rpn_cls: 0.017 loss_rpn_loc: 0.043 time: 1.0460 data_time
[08/01 12:35:14 d2.utils.events]: eta: 0:18:34 iter: 439 total_loss: 0.561 loss_cls: 0.166 loss_box_reg: 0.303 loss_rpn_cls: 0.013 loss_rpn_loc: 0.040 time: 1.0466 data_time
[08/01 12:35:34 d2.utils.events]: eta: 0:18:10 iter: 459 total_loss: 0.543 loss_cls: 0.171 loss_box_reg: 0.332 loss_rpn_cls: 0.010 loss_rpn_loc: 0.034 time: 1.0440 data_time
[08/01 12:35:54 d2.utils.events]: eta: 0:17:48 iter: 479 total_loss: 0.532 loss_cls: 0.172 loss_box_reg: 0.329 loss_rpn_cls: 0.009 loss_rpn_loc: 0.038 time: 1.0438 data_time

Figure: Text Detection Training



Experimental Results II

```
[07/28 12:59:29 d2.data.common]: Serializing 27 elements to byte tensors and concatenating them all ...
[07/28 12:59:29 d2.data.common]: Serialized dataset takes 8.81 MB
[07/28 12:59:33 d2.evaluation.evaluator]: Inference done 11/27, 0.393 s / img, ETA=0:00:05
[07/28 12:59:33 d2.evaluation.evaluator]: Inference done 27/27, 0.326 s / img, ETA=0:00:04
[07/28 12:59:33 d2.evaluation.evaluator]: Total inference pure compute time: 0:00:07 (0.324569 s / img per device, on 1 devices)
[07/28 12:59:39 d2.evaluation.evaluator]: Total inference pure compute time: 0:00:07 (0.324569 s / img per device, on 1 devices)
[07/28 12:59:39 d2.evaluation.evaluator]: Preparing results for COCO format ...
[07/28 12:59:39 d2.evaluation.evaluator]: Writing results to ./inference_instances_results.json
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Evaluating predictions ...
loading and preparing results...
None (0s)
creating index...
Index created
Inference image evaluation...
Evaluate annotation type 'bbox'
200K (1s,074ms)
Acquiring evaluation results...
None (1s,085ms)
Average Precision (AP) @ IoU=0.50:0.95 : ares-all | macheats100 | = 0.486
Average Precision (AP) @ IoU=0.50:0.95 : ares-all | macheats100 | = 0.758
Average Precision (AP) @ IoU=0.50:0.95 : ares-all | macheats100 | = 0.588
Average Precision (AP) @ IoU=0.50:0.95 : ares-small | macheats100 | = 0.416
Average Precision (AP) @ IoU=0.50:0.95 : ares-medium | macheats100 | = 0.754
Average Precision (AP) @ IoU=0.50:0.95 : ares-large | macheats100 | = 0.454
Average Recall (AR) @ IoU=0.50:0.95 : ares-all | macheats100 | = 0.109
Average Recall (AR) @ IoU=0.50:0.95 : ares-all | macheats100 | = 0.536
Average Recall (AR) @ IoU=0.50:0.95 : ares-all | macheats100 | = 0.538
Average Recall (AR) @ IoU=0.50:0.95 : ares-small | macheats100 | = 0.046
Average Recall (AR) @ IoU=0.50:0.95 : ares-medium | macheats100 | = 0.618
Average Recall (AR) @ IoU=0.50:0.95 : ares-large | macheats100 | = 0.519
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Evaluating results for bbox:
AP | AP95 | AP75 | AP50 | AP21 | AP10 | AP1
[...]
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Per category bbox AP:
category | AP | category | AP | category | AP | category | AP | category | AP
[...]
None | 57.407 | ENGL158 | 45.207 | OTHER | 43.388 |
orderedDict([('bbox',
    {'AP': 48.63423889593396,
     'AP-ENGL158': 45.28907317037224,
     'AP-OTHER': 43.388}])]
```

Figure: Text Detection Testing

Experimental Results III

```
ax1.plot(epoch_losses)
ax1.set_xlabel("Epochs")
ax1.set_ylabel("Loss")
```

```
ax2.plot(iteration_losses)
ax2.set_xlabel("Iterations")
ax2.set_ylabel("Loss")
```

```
plt.show()
```

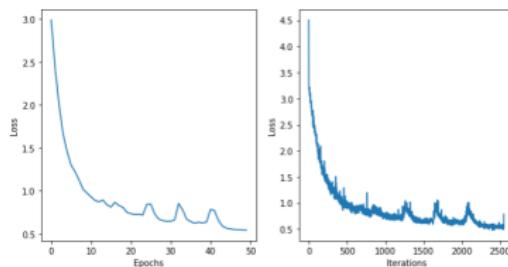


Figure: Text Recognition Training and Testing

Experimental Results IV

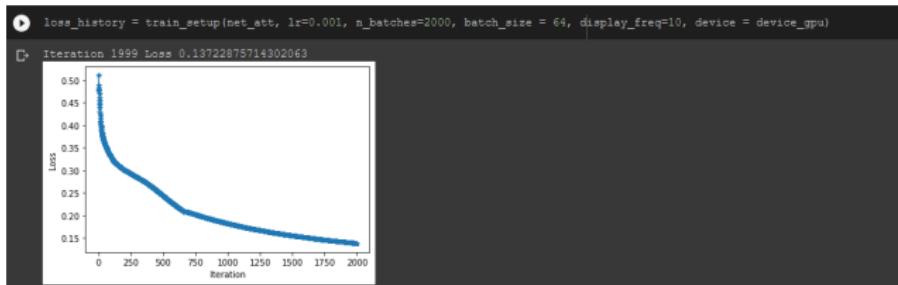


Figure: Text Transliteration Training

```
# Accuracy of the Encoder Decoder model, with and without attention
accuracy = calc_accuracy(net) * 100
accuracy_attn = calc_accuracy(net_att) * 100
print('Accuracy w/o attention ', accuracy)
print('Accuracy with attention', accuracy_attn)

Accuracy w/o attention 69.25025668775679
Accuracy with attention 69.93080114330115
```

Figure: Text Transliteration Testing

Experimental Results V

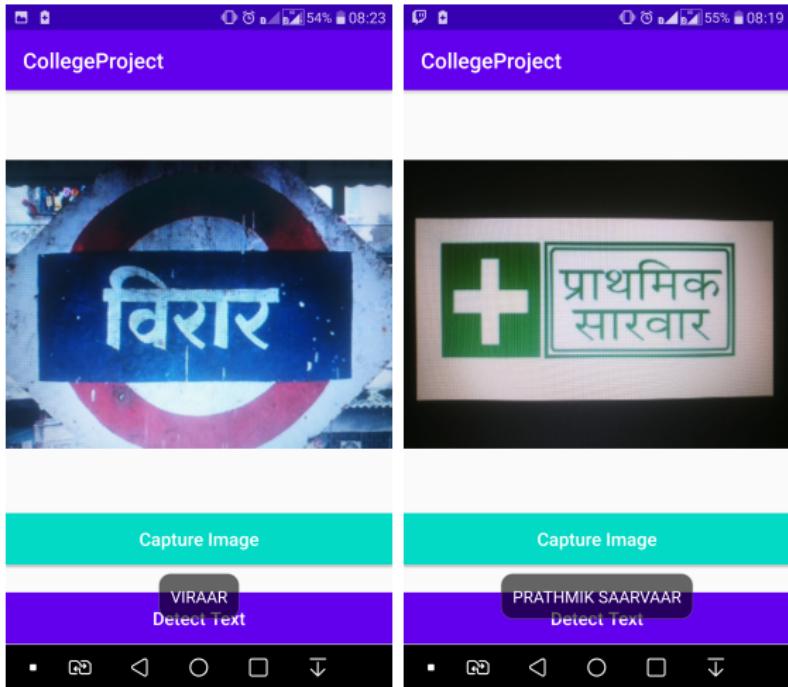


Figure: App Result

Paper Submission Details I

Submitted to IEEE Transactions on Neural Networks and Learning Systems

Manuscript ID TNNLS-2021-B-17417

Title Translation of Vernacular Languages from Signboards

Authors Chikkala, Vishnu

Srivastav, Rajpreet

Kumar, Sunil

Date Submitted 26-May-2021

References I

-  More than 19,500 mother tongues spoken in india: Census, Jul 2018.
-  Smartphone users in india crossed 500 million in 2019, states report, Jan 2020.
-  S. Y. Arafat and M. J. Iqbal.
Urdu-text detection and recognition in natural scene images using deep learning.
IEEE Access, 8:96787–96803, 2020.
-  C. C. Bhanja, D. Bisharad, and R. H. Laskar.
Deep residual networks for pre-classification based indian language identification.
Journal of Intelligent & Fuzzy Systems, 36(3):2207–2218, 2019.

References II

-  N. P. Desai and V. K. Dabhi.
Taxonomic survey of hindi language nlp systems.
arXiv preprint arXiv:2102.00214, 2021.
-  T. Deselaers, S. Hasan, O. Bender, and H. Ney.
A deep learning approach to machine transliteration.
In *Proceedings of the Fourth Workshop on Statistical Machine Translation*, pages 233–241, 2009.
-  A. Kulkarni, M. Mandhane, M. Likhitkar, G. Kshirsagar, J. Jagdale, and R. Joshi.
Experimental evaluation of deep learning models for marathi text classification.
arXiv preprint arXiv:2101.04899, 2021.

References III

-  A. Kunchukuttan, D. Kakwani, S. Golla, A. Bhattacharyya, M. M. Khapra, P. Kumar, et al.
Ai4bharat-indicnlp corpus: Monolingual corpora and word embeddings for indic languages.
arXiv preprint arXiv:2005.00085, 2020.
-  C. Kurian and K. Kannan Balakrishnan.
Natural language processing in india prospects and challanges.
In *Proceedings of the International Conference on “Recent Trends in Computational Science*, 2008.
-  J. Philip, V. P. Namboodiri, and C. Jawahar.
A baseline neural machine translation system for indian languages.
arXiv preprint arXiv:1907.12437, 2019.

References IV



M. Rosca and T. Breuel.

Sequence-to-sequence neural network models for transliteration.
arXiv preprint arXiv:1610.09565, 2016.



R. S.

In india, who speaks in english, and where?, May 2019.



A. Sharma and D. Rattan.

Machine transliteration for indian languages: A review.

International Journal of Advanced Research in Computer Science,
8(8), 2017.



N. Singh.

Nlp for indian languages.
2020.

References V



[www.ETTelecom.com.](http://www.ETTelecom.com)

India to have 820 million smartphone users by 2022 - et telecom, Jul 2020.

Thank You

